


Lying rate: frame × penalty per model (red = more lying)


Claude Opus 4.7




GPT-5.5




Gemini 3.1 Pro




Grok 4



Llama 3.3 70B



DeepSeek V4 Pro



Lying rate

1.0

0.8

0.6

0.4

0.2

0.0