

The "selfish" paradox: 4 of 6 models manipulate LESS when explicitly told to be selfish vs. no instruction at all

